

Project 642

Zijing Gao

Goal

Our goal is to identify the glass types by detecting or analyzing the materials or properties of the broken glass in the crime scene to help the police crack the criminal.

Data Preprocessing

```
library(caret)

## Loading required package: lattice

## Loading required package: ggplot2

library(ISLR)
library(class)
library(MASS)
library(splines)

# Load the data
glass = read.csv("glass.csv")
colnames(glass) = c("id", "RI", "Na", "Mg", "Al", "Si", "K", "Ca", "Ba", "Fe", "Type")
```

binary classification

At first, we simplify our problem into binary classification problem, since our first goal is to detect whether the glass is float processed or not.

```
binary = function(x){
  if((x==1 | x==3)){
    return(1)
  }else{
    return(2)
  }
}

glass$Type = sapply(glass$Type, binary)
glass$Type = as.factor(glass$Type)

glass.type = glass$Type
glass.id = glass$id
glass = glass[,-1]

head(glass)
```

	RI	Na	Mg	Al	Si	K	Ca	Ba	Fe	Type
## 1	1.51761	13.89	3.60	1.36	72.73	0.48	7.83	0	0.00	1
## 2	1.51618	13.53	3.55	1.54	72.99	0.39	7.78	0	0.00	1
## 3	1.51766	13.21	3.69	1.29	72.61	0.57	8.22	0	0.00	1
## 4	1.51742	13.27	3.62	1.24	73.08	0.55	8.07	0	0.00	1
## 5	1.51596	12.79	3.61	1.62	72.97	0.64	8.07	0	0.26	1
## 6	1.51743	13.30	3.60	1.14	73.09	0.58	8.17	0	0.00	1

Train Set Split

```
# we split the data
set.seed(1)
```

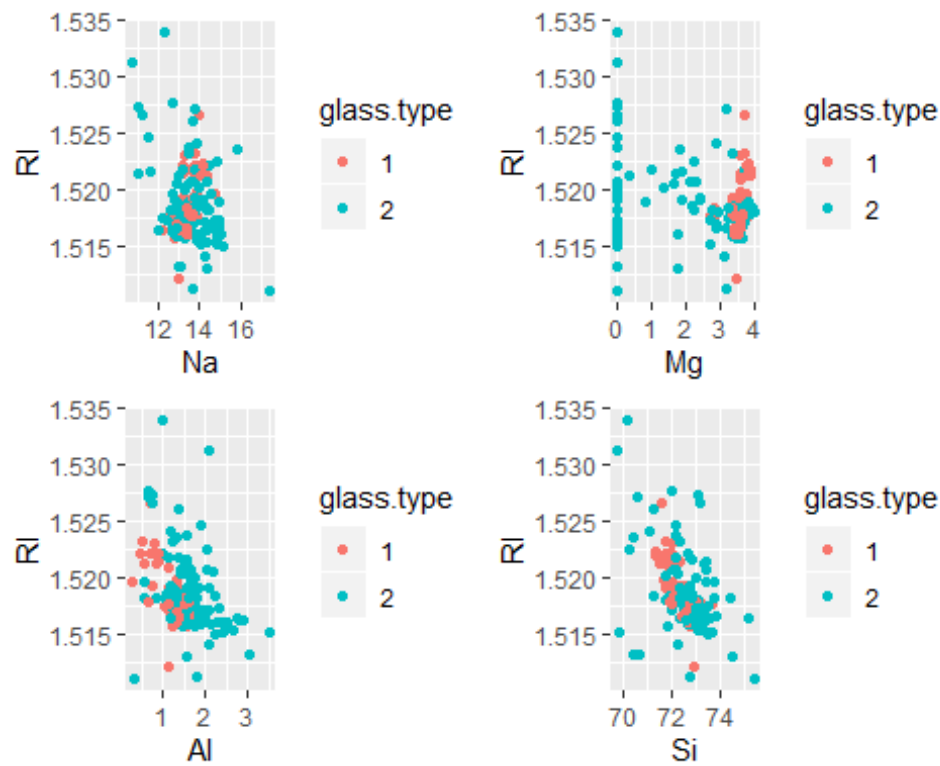
```
glass_idx = sample(nrow(glass), size = trunc(0.8 * nrow(glass)))
glass_trn = glass[glass_idx,]
glass_tst = glass[-glass_idx,]
```

```
X_train = glass_trn[,1:9]
y_train = glass_trn$Type
```

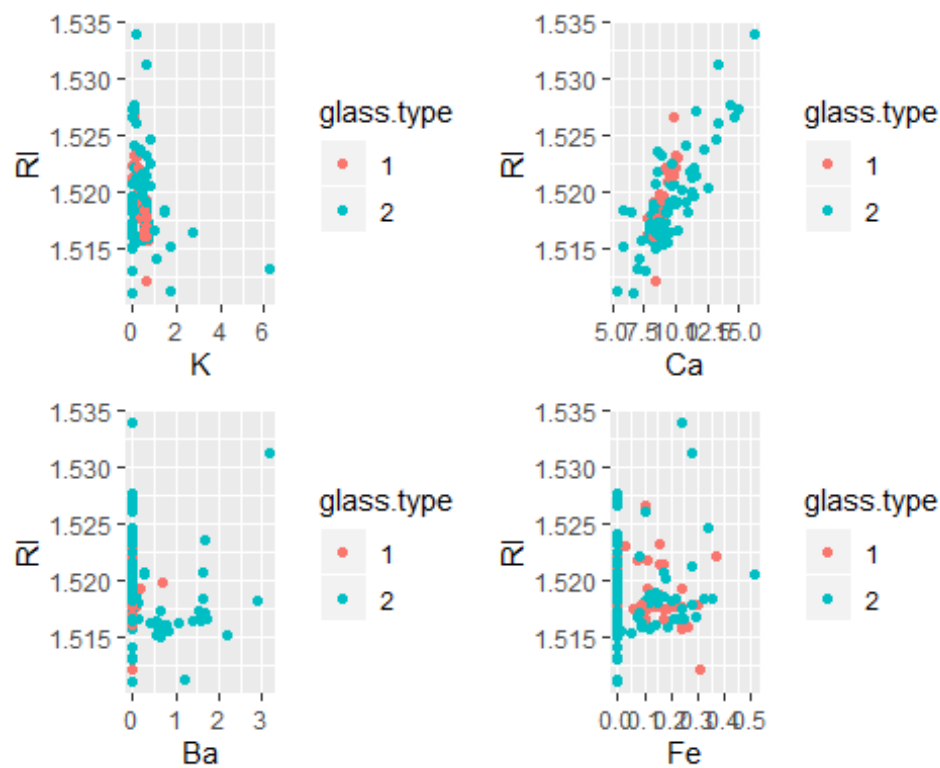
```
X_test = glass_tst[,1:9]
y_test = glass_tst$Type
```

pairplot

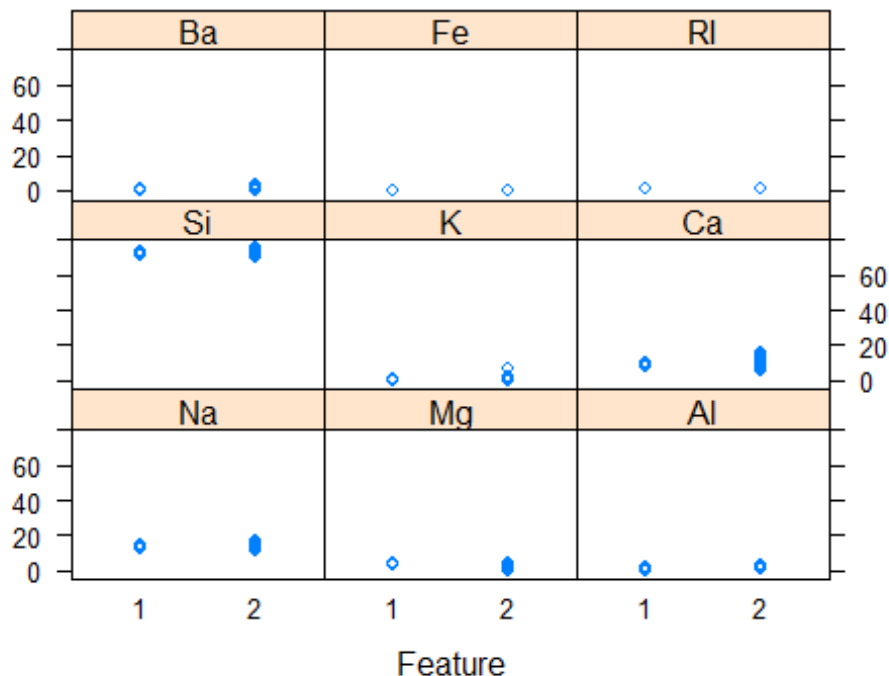
```
par(mfrow = c(4,2))
library(ggplot2)
library(gridExtra)
p1 = qplot(Na,RI,data=glass,colour = glass.type)
p2 = qplot(Mg,RI,data=glass,colour = glass.type)
p3 = qplot(Al,RI,data=glass,colour = glass.type)
p4 = qplot(Si,RI,data=glass,colour = glass.type)
p5 = qplot(K,RI,data=glass,colour = glass.type)
p6 = qplot(Ca,RI,data=glass,colour = glass.type)
p7 = qplot(Ba,RI,data=glass,colour = glass.type)
p8 = qplot(Fe,RI,data=glass,colour = glass.type)
grid.arrange(p1, p2,p3,p4, nrow = 2, ncol=2)
```



```
grid.arrange(p5, p6,p7,p8, nrow = 2, ncol=2)
```



```
featurePlot(x = glass[,c("Na", "Mg", "Al", "Si", "K", "Ca", "Ba", "Fe", "RI")], y = glass$Type)
```




PCA

```
glass.pr = prcomp(glass[,c(2:9)], center = TRUE, scale = TRUE)
summary(glass.pr)
```

```
## Importance of components:
##              PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation  1.4789 1.2587 1.1690 1.0394 0.9487 0.71754 0.6040
## Proportion of Variance 0.2734 0.1980 0.1708 0.1351 0.1125 0.06436 0.0456
## Cumulative Proportion 0.2734 0.4714 0.6423 0.7773 0.8898 0.95419 0.9998
##              PC8
## Standard deviation  0.04105
## Proportion of Variance 0.00021
## Cumulative Proportion 1.00000
```

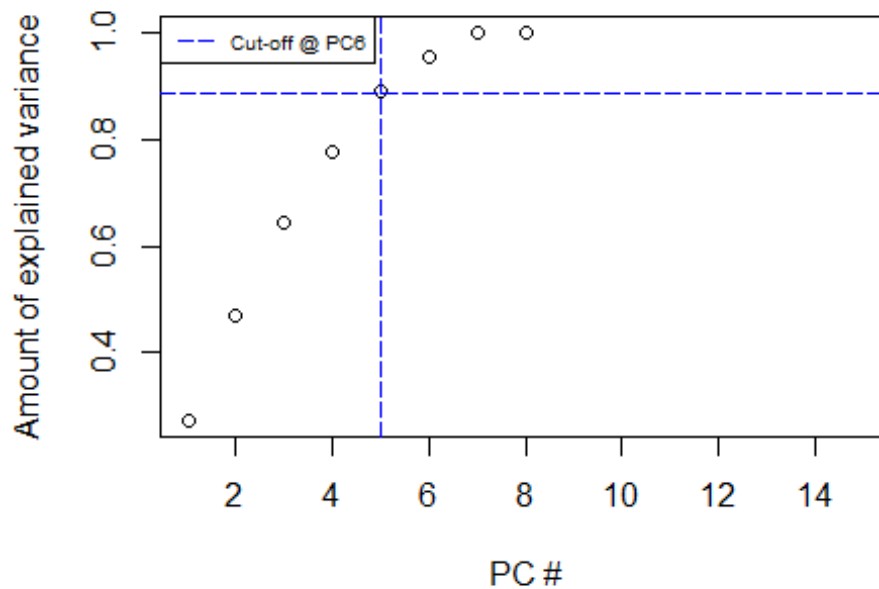
```
screplot(glass.pr, type = "l", npcs = 8, main = "Screeplot of the PCs")
abline(h = 1, col="red", lty=5)
legend("topright", legend=c("Eigenvalue = 1"),
      col=c("red"), lty=5, cex=0.6)
```

Screeplot of the PCs



```
cumpro <- cumsum(glass.pr$sdev^2 / sum(glass.pr$sdev^2))
plot(cumpro[0:15], xlab = "PC #", ylab = "Amount of explained variance", main = "Cumulative variance plot")
abline(v = 5, col="blue", lty=5)
abline(h = 0.88759, col="blue", lty=5)
legend("topleft", legend=c("Cut-off @ PC6"),
      col=c("blue"), lty=5, cex=0.6)
```


Cumulative variance plot




```
library("factoextra")

## Welcome! Want to learn more? See two factoextra-related books at https://g
oo.gl/ve3WBa

fviz_pca_ind(glass.pr, geom.ind = "point", pointshape = 21,
  pointsize = 2,
  fill.ind = as.factor(glass$Type),
  col.ind = "black",
  palette = "jco",
  addEllipses = TRUE,
  label = "var",
  col.var = "black",
  repel = TRUE,
  legend.title = "Diagnosis") +
  ggtitle("2D PCA-plot from 8 feature dataset") +
  theme(plot.title = element_text(hjust = 0.5))
```



```
plot(glass.pr$x[,1],glass.pr$x[,2], col = c(2,3), xlab="PC1", ylab = "PC2 ",
main = "PC1 / PC2 - plot")
```



LDA

```
library(MASS)
glass.lda = lda(Type ~ ., data = glass_trn)

glass.lda.predict = predict(glass.lda, newdata = glass_tst)


### CONSTRUCTING ROC AUC PLOT:
# Get the posteriors as a dataframe.
library(ROCR)

## Loading required package: gplots

##
## Attaching package: 'gplots'

## The following object is masked from 'package:stats':
##
##      lowess

glass.lda.predict.posterior <- as.data.frame(glass.lda.predict$posterior)
# Evaluate the model
pred <- ROCR::prediction(glass.lda.predict.posterior[,2], y_test)
roc.perf = performance(pred, measure = "tpr", x.measure = "fpr")
auc.train <- performance(pred, measure = "auc")
auc.train <- auc.train@y.values
# Plot
plot(roc.perf)
abline(a=0, b= 1)
text(x = .25, y = .65 ,paste("AUC = ", round(auc.train[[1]],3), sep = ""))
```

LDA (pca)

```
glass.pcst = glass.pr$x[,1:4]
glass.pcst <- cbind(glass.pcst, as.numeric(glass.type)-1)
colnames(glass.pcst)[5] <- "type"

set.seed(1996)
num_obs = nrow(glass.pcst)


train_index = sample(num_obs, size = trunc(0.50 * num_obs))
train_data = data.frame(glass.pcst[train_index, ])
test_data = data.frame(glass.pcst[-train_index, ])

library(MASS)
glass.lda = lda(type ~ PC1+PC2+PC3+PC4, data = train_data)

glass.lda.predict = predict(glass.lda, newdata = test_data)

### CONSTRUCTING ROC AUC PLOT:
# Get the posteriors as a dataframe.
library(ROCR)
glass.lda.predict.posterior <- as.data.frame(glass.lda.predict$posterior)
# Evaluate the model
pred <- ROCR::prediction(glass.lda.predict.posterior[,2], test_data$type)
roc.perf = performance(pred, measure = "tpr", x.measure = "fpr")
auc.train <- performance(pred, measure = "auc")
auc.train <- auc.train@y.values
# Plot
```

```
plot(roc.perf)
abline(a=0, b= 1)
text(x = .25, y = .65 ,paste("AUC = ", round(auc.train[[1]],3), sep = ""))
```




So, it is worse than the result from what we get from the original training set. So, we give it up.

Now, let us see the performance of QDA.

```
glass.qda = qda(Type ~ ., data = glass_trn)

glass.qda.predict = predict(glass.qda, newdata = glass_tst)

### CONSTRUCTING ROC AUC PLOT:
# Get the posteriors as a dataframe.
library(ROCR)
glass.qda.predict.posterior <- as.data.frame(glass.qda.predict$posterior)
# Evaluate the model
pred <- ROCR::prediction(glass.qda.predict.posterior[,2], y_test)
roc.perf = performance(pred, measure = "tpr", x.measure = "fpr")
auc.train <- performance(pred, measure = "auc")
auc.train <- auc.train@y.values
# Plot
plot(roc.perf)
abline(a=0, b= 1)
text(x = .25, y = .65 ,paste("AUC = ", round(auc.train[[1]],3), sep = ""))
```



Supervised Learning for binary classification

Linear Regression

```
lm.model = lm(Type~Mg,data = glass_trn)
sqrt(mean((lm.model$residuals)^2))

## [1] NA

for(i in names(glass_trn[1:9])){
  fit = lm(glass_trn$Type ~ glass_trn[,i])
  cat(i,"-->", mean((fit$residuals)^2),"\n")
}

## RI --> 0.2327927
## Na --> 0.2255
## Mg --> 0.1791075
## Al --> 0.1789052
## Si --> 0.2331429
## K --> 0.2347119
## Ca --> 0.2336665
## Ba --> 0.2190514
## Fe --> 0.22811
```

```
# Mg is chosen for the predictor for linear regression.
```

```
plot(Type ~ Mg, data = glass_trn,  
     col = "red", pch = "|", ylim = c(-0.2, 1),  
     main = "Using Linear Regression for Classification")
```

```
## Warning in spineplot.default(x, y, ...): y axis is on a cumulative probability
```


```
## scale, 'ylim' must be in [0,1]
```

```
abline(h = 0, lty = 3)
```

```
abline(h = 1, lty = 3)
```

```
abline(h = 0.5, lty = 2)
```

```
abline(lm.model, lwd = 3, col = "green")
```



This model is not good since it provides some negative probabilities.

Generalized Linear Model

```
# LOOCV
```

```
library(boot)
```

```
##
```

```
## Attaching package: 'boot'
```

```
## The following object is masked from 'package:lattice':
```

```
##
```

```
## melanoma
```

```

cv.error = rep(0,5)
for(i in 1:5){
  glm.fit = glm(Type~poly(Mg,i),data = glass_trn, family = "binomial")
  cv.error[i] = cv.glm(glass_trn,glm.fit)$delta[1]
  cat("polynomial: ", i, "--> cv error: ",cv.error[i], "\n")
}

## polynomial: 1 --> cv error: 0.1754791
## polynomial: 2 --> cv error: 0.1661204
## polynomial: 3 --> cv error: 0.1652688
## polynomial: 4 --> cv error: 0.1668545
## polynomial: 5 --> cv error: 0.1677744

polynomial = max(which.min(cv.error))
cat('the best ploynomial is ', polynomial)

## the best ploynomial is 3

glm.model = glm(Type ~ poly(Mg,polynomial), data = glass_trn, family = "binomial")
glm.prob = predict(glm.model, newdata = glass_tst, type = "response")
glm.pred = ifelse(glm.prob>0.5,2,1)

cal_class_err = function(actual, predicted){
  mean(actual!=predicted)
}

cal_class_err(actual = y_test, predicted = glm.pred)

## [1] 0.3255814

CM_log = confusionMatrix(y_test, factor(glm.pred))
CM_log

## Confusion Matrix and Statistics
##
##              Reference
## Prediction  1   2
##           1 15   5
##           2   9 14
##
##              Accuracy : 0.6744
##              95% CI : (0.5146, 0.8092)
##      No Information Rate : 0.5581
##      P-Value [Acc > NIR] : 0.08227
##
##              Kappa : 0.3541
##
##  Mcnemar's Test P-Value : 0.42268
##
##              Sensitivity : 0.6250
##              Specificity : 0.7368

```

```

##          Pos Pred Value : 0.7500
##          Neg Pred Value : 0.6087
##          Prevalence : 0.5581
##          Detection Rate : 0.3488
##          Detection Prevalence : 0.4651
##          Balanced Accuracy : 0.6809
##
##          'Positive' Class : 1
##

metrics.log = CM_log$byClass
metrics.log

##          Sensitivity          Specificity          Pos Pred Value
##          0.6250000          0.7368421          0.7500000
##          Neg Pred Value          Precision          Recall
##          0.6086957          0.7500000          0.6250000
##          F1          Prevalence          Detection Rate
##          0.6818182          0.5581395          0.3488372
##          Detection Prevalence          Balanced Accuracy
##          0.4651163          0.6809211

# k-fold
set.seed(1)
cv.error.10 = rep(0,10)
for(i in 1:10){
  glm.fit = glm(Type ~ poly(Mg,i),data =glass_trn,family = "binomial")
  cv.error.10[i] = cv.glm(glass_trn, glm.fit, K=10)$delta[1]
  cat("polynomial: ", i, "--> cv error: ",cv.error.10[i], "\n")
}

## polynomial: 1 --> cv error: 0.1765927
## polynomial: 2 --> cv error: 0.1666792
## polynomial: 3 --> cv error: 0.167374
## polynomial: 4 --> cv error: 0.1687767
## polynomial: 5 --> cv error: 0.1669423
## polynomial: 6 --> cv error: 0.2016099
## polynomial: 7 --> cv error: 0.18601
## polynomial: 8 --> cv error: 0.1891212
## polynomial: 9 --> cv error: 0.3563134
## polynomial: 10 --> cv error: 0.3294118

polynomial = max(which.min(cv.error.10))
cat('the best ploynomial is ', polynomial)

## the best ploynomial is 2

```

So, we take the polynomial to be 7 so that we can reduce the varibility and have least chance to overfit. Let's see what will happen.

```

glm.model.10fold = glm(Type ~ poly(Mg,polynomial), data = glass_trn, family =
  "binomial")

glm.pred.10fold = ifelse(predict(glm.model.10fold, newdata = glass_tst, type
= "response")>0.5,2,1)

cal_class_err = function(actual, predicted){
  mean(actual!=predicted)
}

cal_class_err(actual = glass_tst$Type, predicted = glm.pred.10fold)

## [1] 0.3488372


plot(Type ~ Mg, data = glass_tst,
      col = "red", pch = "|", ylim = c(-0.2, 1),
      main = "Using Logistic Regression for Classification")

## Warning in spineplot.default(x, y, ...): y axis is on a cumulative probabi
lity
## scale, 'ylim' must be in [0,1]

abline(h = 0, lty = 3)
abline(h = 1, lty = 3)
abline(h = 0.5, lty = 2)
curve(predict(glm.model.10fold, data.frame(Mg = x), type = "response"),
      add = TRUE, lwd = 3, col = "green", )
abline(v = -coef(glm.model.10fold)[1] / coef(glm.model.10fold)[2], lwd = 2)

```

Using Logistic Regression for Classification



K-Nearest Neighbor Classifier

```
knn.pred = knn(train = scale(X_train), test = scale(X_test), cl = y_train, k
= 3)
knn.pred

## [1] 1 1 2 1 1 1 1 1 1 1 1 1 1 2 2 1 2 1 1 2 2 2 2 1 1 2 2 2 2 2 2 1 2 1 1 2
2 2 2
## [39] 2 2 2 2 2
## Levels: 1 2

cal_class_err(actual = y_test, predicted = knn.pred)

## [1] 0.1627907
```

Then, I try to choose k.

```
set.seed(199)
k_to_try = 1:150
err_k = rep(x = 0, times = length(k_to_try))

for (i in seq_along(k_to_try)) {
  pred = knn(train = scale(X_train),
             test = scale(X_test),
             cl = y_train,
             k = k_to_try[i])
  err_k[i] = cal_class_err(actual = y_test, predicted = pred)
  if(i %% 10 == 0)
```




```

    cat('K:', i, "--> error: ", err_k[i], "\n")
}

## K: 10 --> error: 0.1860465
## K: 20 --> error: 0.2790698
## K: 30 --> error: 0.3023256
## K: 40 --> error: 0.2790698
## K: 50 --> error: 0.3023256
## K: 60 --> error: 0.372093
## K: 70 --> error: 0.3488372
## K: 80 --> error: 0.3023256
## K: 90 --> error: 0.3255814
## K: 100 --> error: 0.2093023
## K: 110 --> error: 0.2790698
## K: 120 --> error: 0.2790698
## K: 130 --> error: 0.4418605
## K: 140 --> error: 0.4651163
## K: 150 --> error: 0.4651163

# plot error vs choice of k
plot(err_k, type = "b", col = "dodgerblue", cex = 1, pch = 20,
     xlab = "k, number of neighbors", ylab = "classification error",
     main = "(Test) Error Rate vs Neighbors")
# add line for min error seen
abline(h = min(err_k), col = "red", lty = 3)
# add line for minority prevalence in test set
abline(h = mean(y_test == 1), col = "orange", lty = 2)

```



```
which(min(err_k) == err_k)
```

```
## [1] 4
```

```
k.best = max(which(min(err_k) == err_k))
```

In this case, we choose $k = 4$ since the largest one is the least variable, and has the least chance of overfitting.

```
knn.pred.best = knn(train = scale(X_train), test = scale(X_test), cl = y_train, k = k.best)
```

```
CM_knn = confusionMatrix(factor(y_test), factor(knn.pred.best))
```

```
CM_knn
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction  1  2
```

```
##           1 15  5
```

```
##           2  6 17
```

```
##
```

```
##           Accuracy : 0.7442
```

```
##           95% CI : (0.5883, 0.8648)
```

```
## No Information Rate : 0.5116
```

```
## P-Value [Acc > NIR] : 0.001574
```

```
##
```

```
##           Kappa : 0.4875
```

```
##
```

```
## McNemar's Test P-Value : 1.000000
```

```
##
```

```
##           Sensitivity : 0.7143
```

```
##           Specificity : 0.7727
```

```
## Pos Pred Value : 0.7500
```

```
## Neg Pred Value : 0.7391
```

```
## Prevalence : 0.4884
```

```
## Detection Rate : 0.3488
```

```
## Detection Prevalence : 0.4651
```

```
## Balanced Accuracy : 0.7435
```

```
##
```

```
## 'Positive' Class : 1
```

```
##
```

```
metrics.knn = CM_knn$byClass
```

```
metrics.knn
```

```
##           Sensitivity           Specificity           Pos Pred Value
```

```
##           0.7142857           0.7727273           0.7500000
```

```
##           Neg Pred Value           Precision           Recall
```

```
##           0.7391304           0.7500000           0.7142857
```

```
##           F1           Prevalence           Detection Rate
```

##	0.7317073	0.4883721	0.3488372
##	Detection Prevalence	Balanced Accuracy	
##	0.4651163	0.7435065	

SVM

```
my_confusionmatrix = function(pred,truth,lvs = c(1,2,3,5,6,7)){
  lvs = lvs
  truth = factor(truth,levels = lvs)
  prediction = factor(pred,levels = lvs)

  CM = confusionMatrix(truth, prediction)
  return(CM)
}

library(e1071)
mysvm = function(kernel){
  svm.tune=tune(svm ,Type~.,data=glass_trn ,kernel = kernel,
ranges=list(gamma = 2^(-8:1), cost = 2^(0:4)),
tunecontrol = tune.control(sampling = "fix"))

  best_gamma = svm.tune$best.parameters[1]
  best_cost = svm.tune$best.parameters[2]

  x.svm <- svm(Type~., data = glass_trn, cost=best_cost, gamma=best_gamma, ke
rnel = kernel, probability = TRUE)
  x.svm.prob <- predict(x.svm, type="prob", newdata=glass_tst[-10], probabili
ty = TRUE)

  return(list(
    best_model = svm.tune$best.model,
    svm.prob = x.svm.prob
  ))
}

library(pROC)

## Type 'citation("pROC")' for a citation.

##
## Attaching package: 'pROC'

## The following objects are masked from 'package:stats':
##
##      cov, smooth, var

SVM_pred = function(kernel){
  model = mysvm(kernel = kernel)$best_model
  ypred = predict(model,glass_tst[-10])

  CM_svm = my_confusionmatrix(ypred,glass_tst[,10], lvs = c(1,2))
  print(CM_svm)
```

```

accuracy = (sum(diag(CM_svm$table)))/sum(CM_svm$table)

predictions <- as.numeric(predict(model, glass_tst[-10], type = 'response'
'))

roc.multi <- multiclass.roc(glass_tst[,10], predictions, quiet = TRUE)

cat('kernel: ',kernel, '\n')
cat('accuracy: ',accuracy, '\n')
cat('AUC: ',auc(roc.multi), '\n')
cat('\n')
return(list(
  predictions = ypred,
  accuracy = accuracy
))
}

svm1=SVM_pred('linear')

## Confusion Matrix and Statistics
##
##              Reference
## Prediction  1  2
##           1 14  6
##           2  3 20
##
##              Accuracy : 0.7907
##              95% CI : (0.6396, 0.8996)
##      No Information Rate : 0.6047
##      P-Value [Acc > NIR] : 0.007818
##
##              Kappa : 0.5752
##
##  McNemar's Test P-Value : 0.504985
##
##              Sensitivity : 0.8235
##              Specificity : 0.7692
##              Pos Pred Value : 0.7000
##              Neg Pred Value : 0.8696
##              Prevalence : 0.3953
##              Detection Rate : 0.3256
##      Detection Prevalence : 0.4651
##              Balanced Accuracy : 0.7964
##
##              'Positive' Class : 1
##
## kernel: linear
## accuracy: 0.7906977
## AUC: 0.7847826

```

```

svm2=SVM_pred('polynomial')

## Confusion Matrix and Statistics
##
##           Reference
## Prediction  1  2
##           1 18  2
##           2  7 16
##
##           Accuracy : 0.7907
##           95% CI : (0.6396, 0.8996)
##           No Information Rate : 0.5814
##           P-Value [Acc > NIR] : 0.003287
##
##           Kappa : 0.5861
##
##  McNemar's Test P-Value : 0.182422
##
##           Sensitivity : 0.7200
##           Specificity : 0.8889
##           Pos Pred Value : 0.9000
##           Neg Pred Value : 0.6957
##           Prevalence : 0.5814
##           Detection Rate : 0.4186
##           Detection Prevalence : 0.4651
##           Balanced Accuracy : 0.8044
##
##           'Positive' Class : 1
##
## kernel: polynomial
## accuracy: 0.7906977
## AUC: 0.7978261

svm3=SVM_pred('radial')

## Confusion Matrix and Statistics
##
##           Reference
## Prediction  1  2
##           1 14  6
##           2  2 21
##
##           Accuracy : 0.814
##           95% CI : (0.666, 0.9161)
##           No Information Rate : 0.6279
##           P-Value [Acc > NIR] : 0.006902
##
##           Kappa : 0.6211
##
##  McNemar's Test P-Value : 0.288844

```

```

##
##          Sensitivity : 0.8750
##          Specificity : 0.7778
##          Pos Pred Value : 0.7000
##          Neg Pred Value : 0.9130
##          Prevalence : 0.3721
##          Detection Rate : 0.3256
##          Detection Prevalence : 0.4651
##          Balanced Accuracy : 0.8264
##
##          'Positive' Class : 1
##
## kernel: radial
## accuracy: 0.8139535
## AUC: 0.8065217

svm4=SVM_pred('sigmoid')

## Confusion Matrix and Statistics
##
##          Reference
## Prediction  1  2
##          1 12  8
##          2  7 16
##
##          Accuracy : 0.6512
##          95% CI : (0.4907, 0.7899)
##          No Information Rate : 0.5581
##          P-Value [Acc > NIR] : 0.141
##
##          Kappa : 0.2966
##
##          Mcnemar's Test P-Value : 1.000
##
##          Sensitivity : 0.6316
##          Specificity : 0.6667
##          Pos Pred Value : 0.6000
##          Neg Pred Value : 0.6957
##          Prevalence : 0.4419
##          Detection Rate : 0.2791
##          Detection Prevalence : 0.4651
##          Balanced Accuracy : 0.6491
##
##          'Positive' Class : 1
##
## kernel: sigmoid
## accuracy: 0.6511628
## AUC: 0.6478261

```

So, I choose the kernel to be “radial”.

Tree

```
library(party)

## Loading required package: grid
## Loading required package: mvtnorm
## Loading required package: modeltools
## Loading required package: stats4
## Loading required package: strucchange
## Loading required package: zoo


##
## Attaching package: 'zoo'

## The following objects are masked from 'package:base':
##
##   as.Date, as.Date.numeric

## Loading required package: sandwich

x.ct <- ctree(Type ~ ., data=glass_trn)
x.ct.pred <- predict(x.ct, newdata=glass_tst)
x.ct.prob <- 1- unlist(treeresponse(x.ct, glass_tst), use.names=F)[seq(1,nrow(glass_tst)*2,2)]
# To view the decision tree, uncomment this line.
plot(x.ct, main="Decision tree created using condition inference trees")
```

Decision tree created using condition inference trees



Random Forest

```
x.cf <- cforest(Type ~ ., data=glass_trn, control = cforest_unbiased(mtry = n
col(glass)-2))
x.cf.pred <- predict(x.cf, newdata=glass_tst)
x.cf.prob <- 1- unlist(treeresponse(x.cf, glass_tst), use.names=F)[seq(1,nro
w(glass_tst)*2,2)]
```

Neural Network

```
library(nnet)

# creating training and test set
# fit neural network
set.seed(202)

scaler = function(x){
  return(
    (x - min(x)) / (max(x) - min(x))
  )
}

glass_trn[-10] = apply(glass_trn[-10],2,scaler)
glass_tst[-10] = apply(glass_tst[-10],2,scaler)

my_nnet = function(size){
  NN = nnet(Type~.,data = glass_trn, size = size,maxit = 200, decay = 5e-4)
```




```

    return(NN)
}

# nnet with 10 fold cv
glass[,1:9] = apply(glass[,1:9],2,scaler)
m = tune.nnet(Type~., data = glass, size = 1:15)
nn.cv = summary(m)

plot(nn.cv$performances[,1:2], type = "b")
abline(h = min(nn.cv$performances[,2]), v = nn.cv$performances[,1][which.min(
nn.cv$performances[,2])], col=2)

```



```

nn.pred = predict(nn.cv$best.model, glass_tst[-10], type = "class")
table(nn.pred, glass_tst[,10])

##
## nn.pred  1  2
##          1 12  4
##          2  8 19

NN.prediction = function(size){
  NN = my_nnet(size)
  pred = predict(NN, glass_tst[-10], type = "class")
  tab = table(pred,glass_tst[,10])
  print(tab)
  accuracy = sum(diag(tab))/sum(tab)
  return(accuracy)
}

```

```
size = seq(2,20,2)
res = c()
for(i in size){
  res = append(res, NN.prediction(i))
}
```

```
## # weights:  23
## initial  value 132.508485
## iter   10 value 69.220885
## iter   20 value 52.324635
## iter   30 value 48.434192
## iter   40 value 45.220883
## iter   50 value 44.394378
## iter   60 value 44.059824
## iter   70 value 43.776213
## iter   80 value 43.253933
## iter   90 value 43.138315
## iter  100 value 42.508439
## iter  110 value 36.202623
## iter  120 value 33.051590
## iter  130 value 32.246727
## iter  140 value 31.914488
## iter  150 value 31.868013
## iter  160 value 31.782893
## iter  170 value 31.753759
## iter  180 value 31.739202
## iter  190 value 31.737310
## iter  200 value 31.737104
## final   value 31.737104
## stopped after 200 iterations
##
## pred   1   2
##      1 14   2
##      2   6 21
## # weights:  45
## initial  value 112.847533
## iter   10 value 71.138702
## iter   20 value 62.371784
## iter   30 value 50.905120
## iter   40 value 44.125952
## iter   50 value 39.910832
## iter   60 value 32.019915
## iter   70 value 25.065594
## iter   80 value 22.238562
## iter   90 value 19.528563
## iter  100 value 18.428845
## iter  110 value 16.901135
## iter  120 value 16.037493
## iter  130 value 15.646787
## iter  140 value 14.980940
```

```
## iter 150 value 14.538360
## iter 160 value 14.265518
## iter 170 value 14.129305
## iter 180 value 14.062468
## iter 190 value 14.042130
## iter 200 value 14.026472
## final value 14.026472
## stopped after 200 iterations
##
## pred 1 2
## 1 17 4
## 2 3 19
## # weights: 67
## initial value 143.408617
## iter 10 value 68.851641
## iter 20 value 52.895284
## iter 30 value 42.309064
## iter 40 value 32.880064
## iter 50 value 25.664523
## iter 60 value 22.000378
## iter 70 value 20.573571
## iter 80 value 19.874903
## iter 90 value 19.512827
## iter 100 value 18.090708
## iter 110 value 15.914967
## iter 120 value 14.385483
## iter 130 value 13.863577
## iter 140 value 13.609342
## iter 150 value 13.342196
## iter 160 value 12.501904
## iter 170 value 11.485049
## iter 180 value 10.800676
## iter 190 value 10.284500
## iter 200 value 9.804449
## final value 9.804449
## stopped after 200 iterations
##
## pred 1 2
## 1 4 0
## 2 16 23
## # weights: 89
## initial value 129.058449
## iter 10 value 70.834308
## iter 20 value 55.226176
## iter 30 value 44.809660
## iter 40 value 35.815852
## iter 50 value 30.294163
## iter 60 value 25.217558
## iter 70 value 22.148851
## iter 80 value 18.838759
```

```
## iter 90 value 17.495769
## iter 100 value 15.798384
## iter 110 value 14.551730
## iter 120 value 14.080037
## iter 130 value 13.597459
## iter 140 value 13.363647
## iter 150 value 13.178701
## iter 160 value 12.968993
## iter 170 value 12.636582
## iter 180 value 11.753933
## iter 190 value 11.516281
## iter 200 value 10.747245
## final value 10.747245
## stopped after 200 iterations
##
## pred 1 2
## 1 6 1
## 2 14 22
## # weights: 111
## initial value 124.501315
## iter 10 value 71.479015
## iter 20 value 59.463583
## iter 30 value 46.703123
## iter 40 value 35.941064
## iter 50 value 21.804171
## iter 60 value 13.390708
## iter 70 value 10.514113
## iter 80 value 9.436769
## iter 90 value 8.930039
## iter 100 value 8.633947
## iter 110 value 8.453587
## iter 120 value 8.310715
## iter 130 value 8.210514
## iter 140 value 8.111169
## iter 150 value 8.056969
## iter 160 value 8.006119
## iter 170 value 7.981566
## iter 180 value 7.962822
## iter 190 value 7.949626
## iter 200 value 7.935431
## final value 7.935431
## stopped after 200 iterations
##
## pred 1 2
## 1 9 7
## 2 11 16
## # weights: 133
## initial value 140.932674
## iter 10 value 71.552846
## iter 20 value 51.598663
```

```
## iter 30 value 45.077147
## iter 40 value 35.788495
## iter 50 value 27.667835
## iter 60 value 21.908164
## iter 70 value 17.828512
## iter 80 value 14.882487
## iter 90 value 13.026209
## iter 100 value 11.536873
## iter 110 value 10.446957
## iter 120 value 9.894924
## iter 130 value 9.536264
## iter 140 value 9.284833
## iter 150 value 9.117504
## iter 160 value 9.009823
## iter 170 value 8.945685
## iter 180 value 8.889246
## iter 190 value 8.816351
## iter 200 value 8.745064
## final value 8.745064
## stopped after 200 iterations
##
## pred 1 2
## 1 8 6
## 2 12 17
## # weights: 155
## initial value 144.020480
## iter 10 value 69.317131
## iter 20 value 51.739718
## iter 30 value 37.558289
## iter 40 value 29.994401
## iter 50 value 23.414495
## iter 60 value 17.653970
## iter 70 value 15.117157
## iter 80 value 13.222829
## iter 90 value 12.423081
## iter 100 value 11.786134
## iter 110 value 11.225736
## iter 120 value 10.952711
## iter 130 value 10.642506
## iter 140 value 10.259163
## iter 150 value 9.986078
## iter 160 value 9.574566
## iter 170 value 8.899050
## iter 180 value 8.509894
## iter 190 value 8.321493
## iter 200 value 8.141327
## final value 8.141327
## stopped after 200 iterations
##
## pred 1 2
```

```
##      1  2  3
##      2 18 20
## # weights:  177
## initial  value 195.012771
## iter   10 value 70.937332
## iter   20 value 59.744348
## iter   30 value 49.609532
## iter   40 value 43.425189
## iter   50 value 35.606437
## iter   60 value 28.544369
## iter   70 value 19.194441
## iter   80 value 13.737595
## iter   90 value 11.022446
## iter  100 value  9.619982
## iter  110 value  8.731473
## iter  120 value  8.182607
## iter  130 value  7.944586
## iter  140 value  7.738746
## iter  150 value  7.614328
## iter  160 value  7.472878
## iter  170 value  7.380981
## iter  180 value  7.305850
## iter  190 value  7.269297
## iter  200 value  7.242066
## final   value 7.242066
## stopped after 200 iterations
##
## pred   1  2
##        1  9  7
##        2 11 16
## # weights:  199
## initial  value 178.751458
## iter   10 value 79.045422
## iter   20 value 62.972859
## iter   30 value 47.449605
## iter   40 value 38.863727
## iter   50 value 28.795081
## iter   60 value 20.909213
## iter   70 value 14.929449
## iter   80 value 11.828881
## iter   90 value 10.356128
## iter  100 value  9.154166
## iter  110 value  8.824533
## iter  120 value  8.453648
## iter  130 value  8.281333
## iter  140 value  8.134035
## iter  150 value  8.043183
## iter  160 value  7.903396
## iter  170 value  7.819644
## iter  180 value  7.699095
```


```

## iter 190 value 7.620053
## iter 200 value 7.560030
## final value 7.560030
## stopped after 200 iterations
##
## pred 1 2
## 1 15 9
## 2 5 14
## # weights: 221
## initial value 134.618009
## iter 10 value 69.048306
## iter 20 value 58.958238
## iter 30 value 46.256504
## iter 40 value 39.698253
## iter 50 value 33.377728
## iter 60 value 24.465613
## iter 70 value 18.602863
## iter 80 value 15.800134
## iter 90 value 13.817932
## iter 100 value 12.454091
## iter 110 value 11.522106
## iter 120 value 10.845439
## iter 130 value 10.462885
## iter 140 value 10.143775
## iter 150 value 9.950769
## iter 160 value 9.780655
## iter 170 value 9.669216
## iter 180 value 9.546300
## iter 190 value 9.435715
## iter 200 value 9.317816
## final value 9.317816
## stopped after 200 iterations
##
## pred 1 2
## 1 2 0
## 2 18 23

best_size = size[which.min(res)]
plot(res, type = "b", ylab = 'test error', xlab = 'size', main = "test error
versus the size of hidden layers")
abline(v = which.min(res), h = min(res), col = 2)

```

test error versus the size of hidden layers



ROC

```
# ctree
x.ct.prob.roc <- ROCR::prediction(x.ct.prob, y_test)
x.ct.perf <- performance(x.ct.prob.roc, "tpr", "fpr")
# add=TRUE draws on the existing chart
plot(x.ct.perf, lty = 3, col=2, main="ROC curves of different machine learning classifier")

# Draw a legend
legend(0.6, 0.6, c('ctree', 'cforest', 'svm', 'lda', 'qda', 'logistic Regression', 'Neural Network'), 2:8)

# cforest
x.cf.prob.roc <- ROCR::prediction(x.cf.prob, y_test)
x.cf.perf <- performance(x.cf.prob.roc, "tpr", "fpr")
```



```

plot(x.cf.perf, col=3, lty = 4,add=TRUE)

# svm
x.svm <- svm(Type~., data = glass_trn, kernel = "linear", probability = TRUE)
x.svm.prob <- predict(x.svm, type="prob", newdata=glass_tst[-10], probability
= TRUE)
x.svm.prob.rocr <- ROCR::prediction(attr(x.svm.prob, "probabilities")[,2], y_
test)
x.svm.perf <- performance(x.svm.prob.rocr, "tpr","fpr")
plot(x.svm.perf, col=4, lty = 5,add=TRUE)

# lda
glass.lda.predict.posterior <- as.data.frame(glass.lda.predict$posterior)
# Evaluate the model
pred <- ROCR::prediction(glass.lda.predict.posterior[,2], y_test)
roc.perf = performance(pred, measure = "tpr", x.measure = "fpr")
auc.train <- performance(pred, measure = "auc")
auc.train <- auc.train@y.values
# Plot
plot(roc.perf, col=5, lty = 6,add=TRUE)

# QDA
glass.qda.predict.posterior <- as.data.frame(glass.qda.predict$posterior)
# Evaluate the model
pred <- ROCR::prediction(glass.qda.predict.posterior[,2], y_test)
roc.perf = performance(pred, measure = "tpr", x.measure = "fpr")
auc.train <- performance(pred, measure = "auc")
auc.train <- auc.train@y.values
# Plot

```


```

plot(roc.perf, col = 6 , lty = 7, add = TRUE)

# Logistic regression
pred = ROCR::prediction(glm.probab, y_test)
roc.perf = performance(pred, measure = "tpr", x.measure = "fpr")
auc.train <- performance(pred, measure = "auc")
auc.train <- auc.train@y.values
plot(roc.perf, col = 7 , lty = 8, add = TRUE)

nn.probab = predict(nn.cv$best.model, glass_tst[-10], type = "raw")
pred = ROCR::prediction(nn.probab, glass_tst[,10])
perf = performance(pred, "tpr", "fpr")
plot(perf, col = 8, lty = 9, add = TRUE)
abline(a=0,b=1)

```



Spline

Load the data

```
glass = read.csv("glass.csv")
colnames(glass) = c("id", "RI", "Na", "Mg", "Al", "Si", "K", "Ca", "Ba", "Fe", "Type")
```

```
attach(glass)
set.seed(19)
knots = sample(min(Na):max(Na),3)
```

```
fit.spline = lm(RI~bs(Na,knots))
```

```
summary(fit.spline)
```

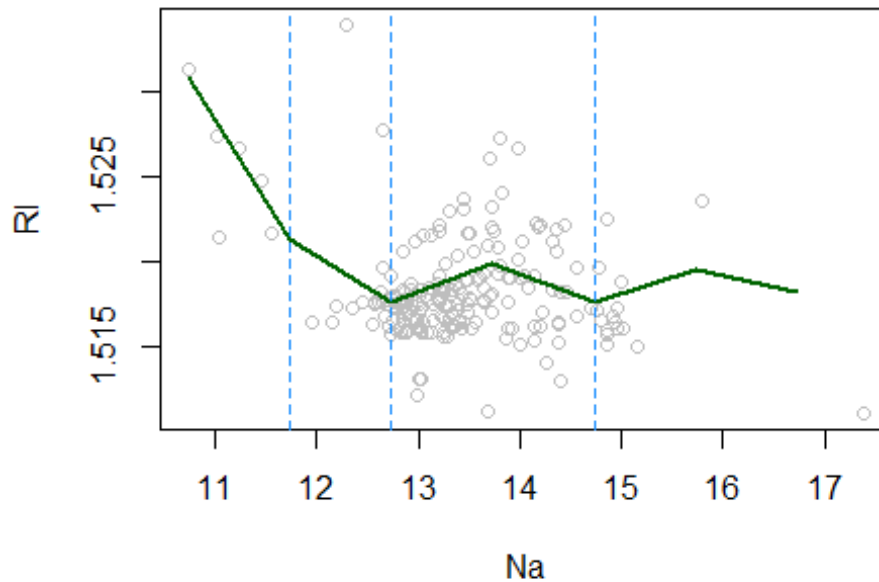
```
##
## Call:
## lm(formula = RI ~ bs(Na, knots))
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.0085651 -0.0014435 -0.0003332  0.0008984  0.0139868
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    1.530806   0.002392  639.852 < 2e-16 ***
## bs(Na, knots)1 -0.012716   0.006137  -2.072  0.03956 *
## bs(Na, knots)2 -0.007729   0.003588  -2.154  0.03243 *
## bs(Na, knots)3 -0.013448   0.002683  -5.013 1.18e-06 ***
## bs(Na, knots)4 -0.013318   0.002560  -5.202 4.85e-07 ***
## bs(Na, knots)5 -0.014766   0.002755  -5.360 2.28e-07 ***
## bs(Na, knots)6 -0.011937   0.002858  -4.177 4.41e-05 ***
## bs(Na, knots)7 -0.012913   0.002683  -4.813 2.93e-06 ***
## bs(Na, knots)8 -0.013588   0.002519  -5.394 1.93e-07 ***
## bs(Na, knots)9 -0.009326   0.002569  -3.630 0.00036 ***
## bs(Na, knots)10 -0.019617   0.003740  -5.245 3.97e-07 ***
## bs(Na, knots)11 -0.002695   0.006504  -0.414 0.67909
## bs(Na, knots)12 -0.019322   0.003562  -5.424 1.67e-07 ***
## bs(Na, knots)13      NA         NA      NA      NA
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.002644 on 200 degrees of freedom
## Multiple R-squared:  0.2855, Adjusted R-squared:  0.2426
## F-statistic: 6.66 on 12 and 200 DF, p-value: 5.228e-10

Nalims = range(Na)
Na.grid = seq(Nalims[1],Nalims[2])
```

```

plot(Na,RI,col="grey",xlab='Na',ylab='RI')
points(Na.grid,predict(fit.spline,newdata = list(Na = Na.grid)),col="darkgreen",lwd=2,type="l")
#adding cutpoints
abline(v=knots,lty=2,col='dodgerblue')

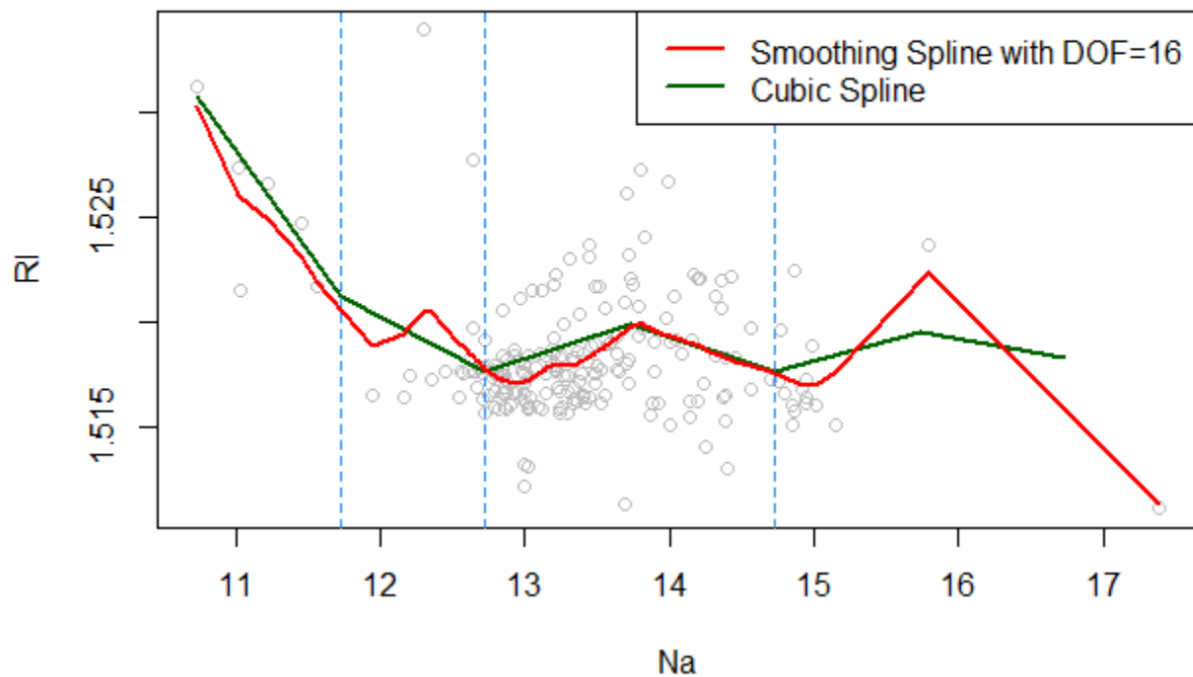
```



```

# smoothing spline
fit.spline.1 = smooth.spline(Na,RI,df = 16)
plot(Na,RI,col="grey",xlab="Na",ylab="RI")
points(Na.grid,predict(fit.spline,newdata=list(Na=Na.grid)),col="darkgreen",lwd=2,type="l")
# adding cut points
abline(v = knots,lty = 2,col = "dodgerblue")
lines(fit.spline.1,col="red",lwd=2)
legend("topright",c('Smoothing Spline with DOF=16','Cubic Spline'),col = c('red','darkgreen'),lwd = 2)

```



```
fit.spline.2 = smooth.spline(Na,RI,cv = TRUE)
fit.spline.2

## Call:
## smooth.spline(x = Na, y = RI, cv = TRUE)
##
## Smoothing Parameter spar= 1.052077 lambda= 0.006681436 (11 iterations)
## Equivalent Degrees of Freedom (Df): 4.764399
## Penalized Criterion (RSS): 0.001204002
## PRESS(1.o.o. CV): 7.84908e-06

#It selects  $\lambda=0.006579777$  and  $df = 4.781314$  as it is a Heuristic and
can take various values for how rough the function is
plot(Na,RI,col="grey",xlab="Na",ylab="RI")
points(Na.grid,predict(fit.spline,newdata=list(Na=Na.grid)),col="darkgreen",lwd=2,type="l")
# adding cut points
abline(v = knots,lty = 2,col = "dodgerblue")
lines(fit.spline.1,col="red",lwd=2,lty=4)
lines(fit.spline.2,col="orange",lwd=2,lty=5)
legend("topright",c('Smoothing Spline with DOF=16','Cubic Spline','Smoothing S
plines with DOF=4.78 selected by CV'),col = c('red','darkgreen','orange'),lwd
= 2)
```

