

## 645 final

Zijing Gao

2019/11/26

```
library('dplyr')

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

library('ggplot2')
library('forecast')

## Registered S3 method overwritten by 'xts':
##   method      from
##   as.zoo.xts zoo

## Registered S3 method overwritten by 'quantmod':
##   method      from
##   as.zoo.data.frame zoo

## Registered S3 methods overwritten by 'forecast':
##   method      from
##   fitted.fracdiff   fracdiff
##   residuals.fracdiff fracdiff

library('tseries')
```

## objective

### plot, examine, and prepare series for modeling

```
# read data
dat = read.csv("weatherHistory.csv", header = T)
```

Now, I need to fetch and manage the data.

At first, I need to remove some columns that we do not need at all.

```

# check colnames
colnames(dat)

## [1] "Formatted.Date"          "Summary"
## [3] "Precip.Type"            "Temperature..C."
## [5] "Apparent.Temperature..C." "Humidity"
## [7] "Wind.Speed..km.h."      "Wind.Bearing..degrees."
## [9] "Visibility..km."        "Loud.Cover"
## [11] "Pressure..millibars."   "Daily.Summary"
## [13] "Precip.Value"

# remove the columns we may not need.
dat.remove = subset(dat, select=-c(Loud.Cover,Daily.Summary,Summary, Precip.Type))

colnames(dat.remove)

## [1] "Formatted.Date"          "Temperature..C."
## [3] "Apparent.Temperature..C." "Humidity"
## [5] "Wind.Speed..km.h."      "Wind.Bearing..degrees."
## [7] "Visibility..km."        "Pressure..millibars."
## [9] "Precip.Value"

# order the data by date.

dat.order = dat.remove[order(as.Date(dat.remove$Formatted.Date, format="%Y-%m-%d")),]
# write.csv(new_data,file = "645_final_data.csv")

# we format the date in a better way so that we can use dplyr to select the subset of the data that we may use.

dat.order$Formatted.Date = as.POSIXct(dat.order$Formatted.Date,format = "%Y-%m-%d %H:%M:%S")

dat.subset = dat.order%>%
  select(Formatted.Date, Temperature..C., Humidity, Wind.Speed..km.h., Wind.Bearing..degrees., Visibility..km., Pressure..millibars.)%>%
  filter(Formatted.Date > "2016-08-01 02:00:00 EST" &
    Formatted.Date < "2016-08-31 02:00:00 EST")

head(dat.subset)

##      Formatted.Date Temperature..C. Humidity Wind.Speed..km.h.
## 1 2016-08-01 03:00:00      19.88333      0.88      19.8513
## 2 2016-08-01 04:00:00      19.81667      0.88      10.6099
## 3 2016-08-01 05:00:00      19.74444      0.89       9.4668
## 4 2016-08-01 06:00:00      18.84444      0.93      12.5580
## 5 2016-08-01 07:00:00      19.93333      0.88       9.5795
## 6 2016-08-01 08:00:00      21.07778      0.87      14.0553

```

```
## Wind.Bearing..degrees. Visibility..km. Pressure..millibars.
## 1 340 15.8263 1012.58
## 2 0 16.1000 1011.40
## 3 57 14.9569 1010.84
## 4 301 9.9015 1011.12
## 5 301 15.8263 1012.64
## 6 313 9.9820 1013.06
```

```
dim(dat.subset)
```

```
## [1] 719 7
```

## use correlation matrix to find explanatory variables

```
dat.cor = subset(dat.remove, select=-c(Formatted.Date))
```

```
corr = cor(dat.cor)
```

```
corr
```


```
## Temperature..C. Apparent.Temperature..C.
## Temperature..C. 1.0000000000 0.9926285642
## Apparent.Temperature..C. 0.9926285642 1.0000000000
## Humidity -0.6322546750 -0.6025709956
## Wind.Speed..km.h. 0.0089569683 -0.0566496983
## Wind.Bearing..degrees. 0.0299882045 0.0290305198
## Visibility..km. 0.3928465717 0.3817184705
## Pressure..millibars. -0.0054471062 -0.0002189998
## Precip.Value 0.0008619473 0.0010608217
## Humidity Wind.Speed..km.h.
## Temperature..C. -0.6322546750 0.0089569683
## Apparent.Temperature..C. -0.6025709956 -0.0566496983
## Humidity 1.0000000000 -0.224951456
## Wind.Speed..km.h. -0.2249514559 1.0000000000
## Wind.Bearing..degrees. 0.0007346454 0.103821508
## Visibility..km. -0.3691725006 0.100749284
## Pressure..millibars. 0.0054542633 -0.049262806
## Precip.Value 0.0002169395 0.004804537
## Wind.Bearing..degrees. Visibility..km.
## Temperature..C. 0.0299882045 0.392846572
## Apparent.Temperature..C. 0.0290305198 0.381718470
## Humidity 0.0007346454 -0.369172501
## Wind.Speed..km.h. 0.1038215077 0.100749284
## Wind.Bearing..degrees. 1.0000000000 0.047594175
## Visibility..km. 0.0475941753 1.0000000000
## Pressure..millibars. -0.0116508848 0.059818381
## Precip.Value -0.0039252047 0.008057509
## Pressure..millibars. Precip.Value
## Temperature..C. -0.0054471062 0.0008619473
## Apparent.Temperature..C. -0.0002189998 0.0010608217
## Humidity 0.0054542633 0.0002169395
## Wind.Speed..km.h. -0.0492628055 0.0048045372
```

```
## Wind.Bearing..degrees.      -0.0116508848 -0.0039252047
## Visibility..km.             0.0598183810  0.0080575088
## Pressure..millibars.        1.0000000000 -0.0096836303
## Precip.Value                -0.0096836303  1.0000000000
```

```
library(corrplot)
```

```
## corrplot 0.84 loaded
```

```
corr = cor(dat.cor)
corrplot(corr, type = "upper", order = "hclust",
         tl.col = "black", tl.srt = 45)
```



```
corr
```

```
##              Temperature..C. Apparent.Temperature..C.
## Temperature..C.      1.0000000000      0.9926285642
## Apparent.Temperature..C. 0.9926285642      1.0000000000
## Humidity              -0.6322546750     -0.6025709956
## Wind.Speed..km.h.      0.0089569683     -0.0566496983
## Wind.Bearing..degrees. 0.0299882045      0.0290305198
## Visibility..km.        0.3928465717      0.3817184705
## Pressure..millibars.   -0.0054471062     -0.0002189998
## Precip.Value          0.0008619473      0.0010608217
##              Humidity Wind.Speed..km.h.
## Temperature..C.      -0.6322546750      0.008956968
## Apparent.Temperature..C. -0.6025709956     -0.056649698
## Humidity              1.0000000000     -0.224951456
```


```
## Wind.Speed..km.h.      -0.2249514559      1.0000000000
## Wind.Bearing..degrees.  0.0007346454      0.103821508
## Visibility..km.        -0.3691725006      0.100749284
## Pressure..millibars.    0.0054542633     -0.049262806
## Precip.Value           0.0002169395      0.004804537
##                          Wind.Bearing..degrees.  Visibility..km.
## Temperature..C.         0.0299882045      0.392846572
## Apparent.Temperature..C. 0.0290305198      0.381718470
## Humidity                 0.0007346454     -0.369172501
## Wind.Speed..km.h.        0.1038215077      0.100749284
## Wind.Bearing..degrees.    1.0000000000      0.047594175
## Visibility..km.          0.0475941753      1.0000000000
## Pressure..millibars.     -0.0116508848      0.059818381
## Precip.Value            -0.0039252047      0.008057509
##                          Pressure..millibars.  Precip.Value
## Temperature..C.         -0.0054471062     0.0008619473
## Apparent.Temperature..C. -0.0002189998     0.0010608217
## Humidity                 0.0054542633     0.0002169395
## Wind.Speed..km.h.        -0.0492628055     0.0048045372
## Wind.Bearing..degrees.   -0.0116508848    -0.0039252047
## Visibility..km.          0.0598183810     0.0080575088
## Pressure..millibars.     1.0000000000    -0.0096836303
## Precip.Value            -0.0096836303     1.0000000000
```

*# choose temperature and visibility*

#Humidity, Temp, Visibility


```
acf(dat.subset$Humidity, main = "ACF for Humidity", lag = 200)
```

**ACF for Humidity**




```
plot.ts(dat.subset$Humidity, main = "Humidity plot")
```

**Humidity plot**




```
acf(dat.subset$Temperature..C., main = "ACF for Temp", lag = 200)
```

**ACF for Temp**




```
plot.ts(dat.subset$Temperature..C., main = "Temp plot")
```

**Temp plot**




```
acf(dat.subset$Wind.Speed..km.h., main = "ACF for Wind Speed", lag = 200)
```

### ACF for Wind Speed



```
plot.ts(dat.subset$Wind.Speed..km.h., main = "Wind Speed plot")
```

### Wind Speed plot






## How to determine order of differencing needed?

### step 2 Examine the data

*# plot the humidity after cleaning the abnormal data if needed.*

```
ggplot(dat.subset, aes(Formatted.Date, Humidity)) + geom_line() + xlab("Date")  
+ ylab("Humidity of Aug 2016")
```




```
count_ts = ts(dat.subset[, c('Humidity')])
```

```
dat.subset$clean_humidity = tsclean(count_ts)
```

```
ggplot() +  
  geom_line(data = dat.subset, aes(x = Formatted.Date, y = clean_humidity))
```

## Don't know how to automatically pick scale for object of type ts. Defaulting to continuous.



## Smooth the data


Visually, we could draw a line through the series tracing its bigger troughs and peaks while smoothing out noisy fluctuations. This line can be described by one of the simplest — but also very useful — concepts in time series analysis known as a moving average.

```
dat.subset$humid_ma_w = ma(dat.subset$clean_humidity, order=7*24) # using the
  clean count with no outliers
dat.subset$humid_ma_d = ma(dat.subset$clean_humidity, order=24)
```

```
ggplot() +
  geom_line(data = dat.subset, aes(x = Formatted.Date, y = clean_humidity, colour = "Counts")) +
  geom_line(data = dat.subset, aes(x = Formatted.Date, y = humid_ma_w, colour = "weekly Moving Average")) +
  geom_line(data = dat.subset, aes(x = Formatted.Date, y = humid_ma_d, colour = "daily Moving Average")) +
  ylab('Humidity') +
  xlab('date')
```

```
## Don't know how to automatically pick scale for object of type ts. Defaulting to continuous.
```

```
## Warning: Removed 168 rows containing missing values (geom_path).  
## Warning: Removed 24 rows containing missing values (geom_path).
```



### step 3: Decompose the data

```
# Decompose the data
```


```
## Does the series appear to have trends or seasonality?
```

```
hum.ts = ts(dat.subset$Humidity, start = c(2016, 8), frequency = 24*7)
```

```
d = decompose(hum.ts)
```


```
plot(d)
```

## Decomposition of additive time series



```
temp.ts = ts(dat.subset$Temperature..C.,start = c(2016,8),frequency = 24*7)
d.temp = decompose(temp.ts)
plot(d.temp)
```


## Decomposition of additive time series



```
wind.ts = ts(dat.subset$Wind.Speed..km.h.,start = c(2016,8),frequency = 24*7)


d.wind = decompose(wind.ts)
plot(d.wind)
```

## Decomposition of additive time series




## relationships

```
fit.temp = lm(dat.subset$Humidity ~ dat.subset$Temperature..C.)  
plot(y = dat.subset$Humidity, x = dat.subset$Temperature..C., type = "p")  
abline(fit.temp, col = "red")
```



```
fit.wind = lm(dat.subset$Humidity ~ dat.subset$Wind.Speed..km.h.)
plot(y = dat.subset$Humidity, x = dat.subset$Wind.Speed..km.h., type = "p")
abline(fit.wind, col = "red")
```



## regression models

```
library(lmtest)
```

```
## Loading required package: zoo
```

```
##
```

```
## Attaching package: 'zoo'
```

```
## The following objects are masked from 'package:base':
```

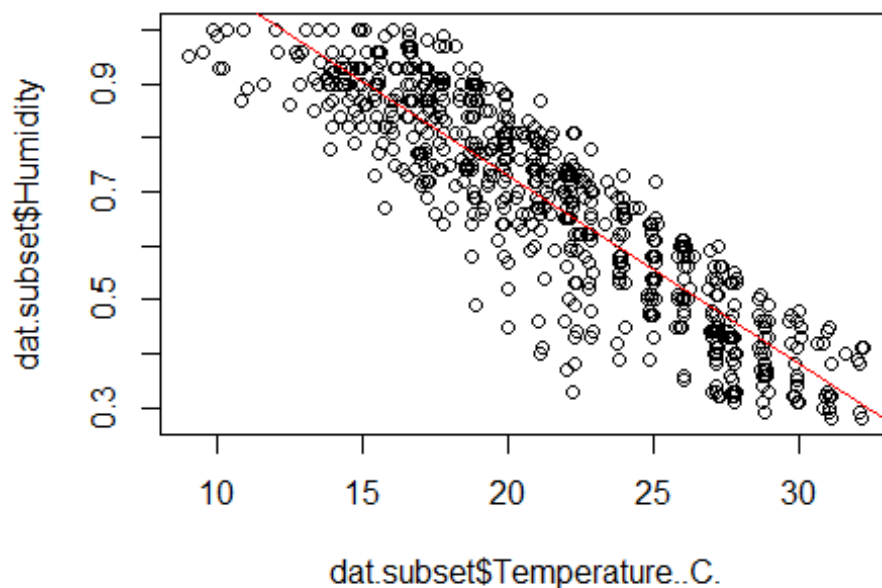
```
##
```

```
##      as.Date, as.Date.numeric
```

```
fit.temp = lm(dat.subset$Humidity ~ dat.subset$Temperature..C.)
```

```
plot(y = dat.subset$Humidity, x = dat.subset$Temperature..C., type = "p")
```

```
abline(fit.temp, col = "red")
```



```
summary(fit.temp)
```

```
##
```

```
## Call:
```

```
## lm(formula = dat.subset$Humidity ~ dat.subset$Temperature..C.)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
```

```
## -0.32253 -0.05610  0.00483  0.07141  0.18057
```


```
##
```

```
## Coefficients:
```



```
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept)      1.4292855   0.0137976  103.59  <2e-16 ***
## dat.subset$Temperature..C. -0.0349542   0.0006258  -55.85  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.08626 on 717 degrees of freedom
## Multiple R-squared:  0.8131, Adjusted R-squared:  0.8129
## F-statistic: 3120 on 1 and 717 DF,  p-value: < 2.2e-16

plot(rstudent(fit.temp))
```



```
dwtest(fit.temp) # < dL: reject H0

##
## Durbin-Watson test
##
## data: fit.temp
## DW = 0.13645, p-value < 2.2e-16
## alternative hypothesis: true autocorrelation is greater than 0

res.temp = fit.temp$residuals
rnum.temp = 0
rdenom.temp = 0
for (i in 2:719){
  rnum.temp = rnum.temp + res.temp[i] * res.temp[i-1]
  rdenom.temp = rdenom.temp + res.temp[i-1]^2
}
```

```

}
rhat.temp = rnum.temp / rdenom.temp
x.temp = rep(0, 718)
y.temp = rep(0, 718)
for (i in 1: 718){
  y.temp[i] = dat.subset$Humidity[i+1] - rhat.temp * dat.subset$Humidity[i]
  x.temp[i] = dat.subset$Temperature..C.[i+1] - rhat.temp * dat.subset$Temperature..C.[i]
}
trans.temp = lm(y.temp ~ x.temp)
summary(trans.temp)


##
## Call:
## lm(formula = y.temp ~ x.temp)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.12066 -0.01798 -0.00032  0.01674  0.15744
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.1021176   0.0016646   61.35  <2e-16 ***
## x.temp       -0.0378432   0.0008116  -46.63  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.03106 on 716 degrees of freedom
## Multiple R-squared:  0.7523, Adjusted R-squared:  0.7519
## F-statistic: 2174 on 1 and 716 DF, p-value: < 2.2e-16

dwtest(trans.temp) # past DW test

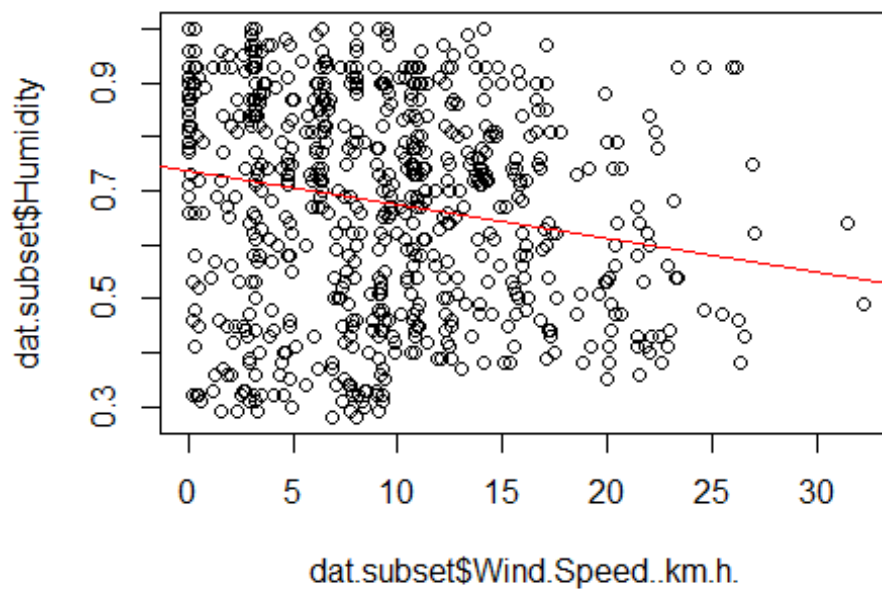
##
## Durbin-Watson test
##
## data: trans.temp
## DW = 2.119, p-value = 0.942
## alternative hypothesis: true autocorrelation is greater than 0

plot(rstudent(trans.temp))

```



```
fit.wind = lm(dat.subset$Humidity ~ dat.subset$Wind.Speed..km.h.)  
plot(y = dat.subset$Humidity, x = dat.subset$Wind.Speed..km.h., type = "p")  
abline(fit.wind, col = "red")
```




```

fit.multi = lm(dat.subset$Humidity ~ dat.subset$Temperature..C.+dat.subset$Wind.Speed..km.h.)
summary(fit.multi)

##
## Call:
## lm(formula = dat.subset$Humidity ~ dat.subset$Temperature..C. +
##     dat.subset$Wind.Speed..km.h.)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.32181 -0.05524  0.00394  0.07045  0.17935
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    1.426926   0.0138635  102.932   <2e-16 ***
## dat.subset$Temperature..C. -0.0351949   0.0006443  -54.629   <2e-16 ***
## dat.subset$Wind.Speed..km.h.  0.0008231   0.0005315    1.549    0.122
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.08618 on 716 degrees of freedom
## Multiple R-squared:  0.8137, Adjusted R-squared:  0.8132
## F-statistic: 1564 on 2 and 716 DF,  p-value: < 2.2e-16

plot(rstudent(fit.multi))

```



```

res.multi = fit.multi$residuals
rnum.multi = 0
rdenom.multi = 0
for (i in 2: 719){
  rnum.multi = rnum.multi + res.multi[i] * res.multi[i-1]
  rdenom.multi = rdenom.multi + res.multi[i-1]^2
}
rhat.multi = rnum.multi / rdenom.multi
x.multi = rep(0, 718)
z.multi = rep(0, 718)
y.multi = rep(0, 718)
for (i in 1: 718){
  y.multi[i] = dat.subset$Humidity[i+1] - rhat.multi * dat.subset$Humidity[i]
  z.multi[i] = dat.subset$Wind.Speed..km.h.[i+1] - rhat.multi * dat.subset$Wind.Speed..km.h.[i]
  x.multi[i] = dat.subset$Temperature..C.[i+1] - rhat.multi * dat.subset$Temperature..C.[i]
}
trans.multi = lm(y.multi ~ x.multi + z.multi)
summary(trans.multi)


##
## Call:
## lm(formula = y.multi ~ x.multi + z.multi)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.120251 -0.017728 -0.000327  0.016950  0.153022
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.1044699   0.0016796   62.198  <2e-16 ***
## x.multi      -0.0376190   0.0008147  -46.174  <2e-16 ***
## z.multi      -0.0007729   0.0003483   -2.219   0.0268 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.03097 on 715 degrees of freedom
## Multiple R-squared:  0.7541, Adjusted R-squared:  0.7534
## F-statistic: 1096 on 2 and 715 DF,  p-value: < 2.2e-16

dwtest(trans.multi) # past DW test

##
## Durbin-Watson test
##
## data:  trans.multi
## DW = 2.1182, p-value = 0.9417
## alternative hypothesis: true autocorrelation is greater than 0

plot(rstudent(trans.multi))

```



```
library(ISLR)
fit.poly = lm(dat.subset$Humidity ~ poly(dat.subset$Temperature..C., degree =
3) + poly(dat.subset$Wind.Speed..km.h., degree = 2))
summary(fit.poly)
```

```
##
## Call:
## lm(formula = dat.subset$Humidity ~ poly(dat.subset$Temperature..C.,
##     degree = 3) + poly(dat.subset$Wind.Speed..km.h., degree = 2))
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-0.32331	-0.05084	0.00533	0.06992	0.17757

```
##
## Coefficients:
```


	Estimate	Std. Error
(Intercept)	0.679889	0.003095
poly(dat.subset\$Temperature..C., degree = 3)1	-4.847601	0.085613
poly(dat.subset\$Temperature..C., degree = 3)2	-0.340298	0.084281
poly(dat.subset\$Temperature..C., degree = 3)3	0.533649	0.083312
poly(dat.subset\$Wind.Speed..km.h., degree = 2)1	0.130584	0.086761
poly(dat.subset\$Wind.Speed..km.h., degree = 2)2	0.096205	0.083420

```
##
## t value Pr(>|t|)
```

	t value	Pr(> t )	
(Intercept)	219.671	< 2e-16	***
poly(dat.subset\$Temperature..C., degree = 3)1	-56.622	< 2e-16	***
poly(dat.subset\$Temperature..C., degree = 3)2	-4.038	5.98e-05	***
poly(dat.subset\$Temperature..C., degree = 3)3	6.405	2.72e-10	***

```
## poly(dat.subset$Wind.Speed..km.h., degree = 2)1 1.505 0.133
## poly(dat.subset$Wind.Speed..km.h., degree = 2)2 1.153 0.249
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.08299 on 713 degrees of freedom
## Multiple R-squared: 0.828, Adjusted R-squared: 0.8268
## F-statistic: 686.4 on 5 and 713 DF, p-value: < 2.2e-16

plot(rstudent(fit.poly))
```




```
fit.poly2 = lm(dat.subset$Humidity ~ poly(dat.subset$Temperature..C., degree
= 3))
summary(fit.poly2)

##
## Call:
## lm(formula = dat.subset$Humidity ~ poly(dat.subset$Temperature..C.,
##     degree = 3))
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.32698 -0.05161  0.00645  0.06897  0.18020
##
## Coefficients:
##                                     Estimate Std. Error t value
## (Intercept)                        0.679889   0.003099  219.415
```


```
## poly(dat.subset$Temperature..C., degree = 3)1 -4.817926 0.083087 -57.986
## poly(dat.subset$Temperature..C., degree = 3)2 -0.349889 0.083087 -4.211
## poly(dat.subset$Temperature..C., degree = 3)3 0.526063 0.083087 6.331
##                                     Pr(>|t|)
## (Intercept)                        < 2e-16 ***
## poly(dat.subset$Temperature..C., degree = 3)1 < 2e-16 ***
## poly(dat.subset$Temperature..C., degree = 3)2 2.87e-05 ***
## poly(dat.subset$Temperature..C., degree = 3)3 4.29e-10 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.08309 on 715 degrees of freedom
## Multiple R-squared:  0.8271, Adjusted R-squared:  0.8264
## F-statistic: 1140 on 3 and 715 DF, p-value: < 2.2e-16

plot(rstudent(fit.poly2))
```



```
count_humid_ma = ts(na.omit(dat.subset$humid_ma_d), frequency=30)
decomp = stl(count_humid_ma, s.window="periodic")
deseasonal_humid <- seasadj(decomp)
plot(decomp)
```





*# we are specifying periodicity of the data, i.e., number of observations per period. Since we are using smoothed daily data, we have 30 observations per month.*

*# We now have a de-seasonalized series*

## step 4: Stationary

### Is the series stationary?


```
adf.test(count_humid_ma, alternative = "stationary")

##
## Augmented Dickey-Fuller Test
##
## data: count_humid_ma
## Dickey-Fuller = -3.658, Lag order = 8, p-value = 0.02711
## alternative hypothesis: stationary
```

## step 5: Autocorrelations and Choosing Model Order

*# We can start with the order of  $d = 1$  and re-evaluate whether further differencing is needed.*

```
count_d1 = diff(deseasonal_humid, differences = 1)
plot(count_d1)
```




```
adf.test(count_d1, alternative = "stationary")  
  
## Warning in adf.test(count_d1, alternative = "stationary"): p-value smaller  
## than printed p-value  
  
##  
## Augmented Dickey-Fuller Test  
##  
## data: count_d1  
## Dickey-Fuller = -5.7047, Lag order = 8, p-value = 0.01  
## alternative hypothesis: stationary
```

## ACF and PACF


```
acf(count_d1, main='ACF for Differenced Series')
```

### ACF for Differenced Series



```
pacf(count_d1, main='PACF for Differenced Series')
```

### PACF for Differenced Series



*# There are significant and constantly decreasing auto correlations within one day from our hourly data and beyond.*

*# Partial correlation plots show a significant spike at lag 1 to 4.*

## Fitting an ARIMA model

*# We can specify non-seasonal ARIMA structure and fit the model to de-seasonalized data. Parameters (1,1,2) suggested by the automated procedure are in line with our expectations based on the steps above*

```
auto.arima(deseasonal_humid, seasonal = FALSE)
```

```
## Series: deseasonal_humid
```

```
## ARIMA(1,1,2)
```

```
##
```

```
## Coefficients:
```

```
##          ar1      ma1      ma2
```

```
##          0.9102  0.8232  0.0940
```

```
## s.e.    0.0170  0.0416  0.0412
```

```
##
```

```
## sigma^2 estimated as 1.836e-06:  log likelihood=3598.06
```

```
## AIC=-7188.12  AICc=-7188.06  BIC=-7169.95
```

Using the ARIMA notation introduced above, the fitted model can be written as

$$Y_{d_t} = 0.9102Y_{t-1} + 0.8232e_{t-1} + 0.0940e_{t-2} + \epsilon$$


where  $\epsilon$  is some error and the original series is differenced with order 1.

AR(1) coefficient  $p = 0.9102$  tells us that the next value in the series is taken as a dampened previous value by a factor of 0.91 and depends on previous error lag.

## step 7: Evaluate and Iterate

```
fit<-auto.arima(deseasonal_humid, seasonal=FALSE)
```

```
tsdisplay(residuals(fit), lag.max=45, main='(1,1,2) Model Residuals')
```




*# There is a clear pattern present in ACF/PACF and model residuals plots repeating at lag 24. This suggests that our model may be better off with a different specification, such as  $p = 24$  or  $q = 24$ .*

```
fit2 = arima(deseasonal_humid, order=c(1,1,24))
fit2


##
## Call:
## arima(x = deseasonal_humid, order = c(1, 1, 24))
##
## Coefficients:
##      ar1      ma1      ma2      ma3      ma4      ma5      ma6      ma7
##    -0.0291  1.7521  1.7750  1.7855  1.8096  1.7269  1.7984  1.8052
## s.e.   0.0853  0.0793  0.1274  0.1480  0.1592  0.1598  0.1485  0.1530
##      ma8      ma9     ma10     ma11     ma12     ma13     ma14     ma15
##    1.7254  1.6384  1.5578  1.4055  1.4101  1.3843  1.2958  1.2245
## s.e.   0.1517  0.1477  0.1486  0.1622  0.1901  0.2214  0.2512  0.2744
##      ma16     ma17     ma18     ma19     ma20     ma21     ma22     ma23
##    1.1809  1.1131  1.2015  1.2949  1.3277  1.3668  1.4518  1.4670
## s.e.   0.2841  0.2825  0.2725  0.2461  0.2147  0.1776  0.1308  0.0859
##      ma24
##    0.6333
## s.e.   0.0487
##
## sigma^2 estimated as 1.091e-06:  log likelihood = 3749.87,  aic = -7447.75
```

```
tsdisplay(residuals(fit2), lag.max=15, main='(1,1,24) Seasonal Model Residuals')
```



```
fcast <- forecast(fit2, h=24)  
plot(fcast)
```

## Forecasts from ARIMA(1,1,24)



### holt-out forecast

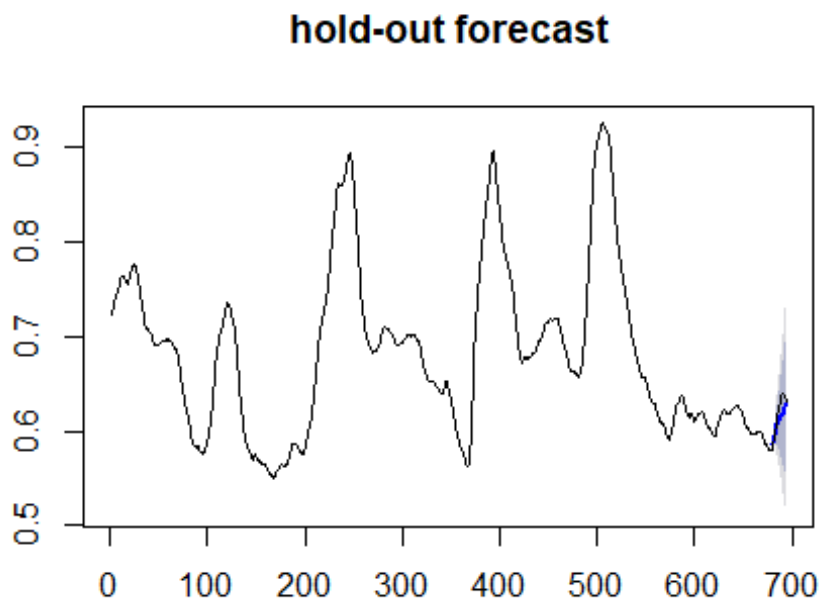
The light blue line above shows the fit provided by the model, but what if we wanted to get a sense of how the model will perform in the future? One method is to reserve a portion of our data as a “hold-out” set, fit the model, and then compare the forecast to the actual observed values:

```
hold <- window(ts(deseasonal_humid), start=680)

fit_no_holdout = arima(ts(deseasonal_humid[-c(680:695)]), order=c(1,1,24))

fcast_no_holdout <- forecast(fit_no_holdout,h=15)

plot(fcast_no_holdout, main="hold-out forecast")
lines(ts(deseasonal_humid))
```




**Testing the distribution of errors in your ARIMA model.**

**Are successive errors correlated?**

```
acf(as.numeric(fcast_no_holdout$residuals), lag.max=20)
```



**Series as.numeric(fcast\_no\_holdout\$residuals)**



```
fit_w_seasonality = auto.arima(deseasonal_humid, seasonal=TRUE)
```

```
fit_w_seasonality
```

```
## Series: deseasonal_humid
```

```
## ARIMA(1,1,2)(0,0,1)[30]
```

```
##
```

```
## Coefficients:
```

```
##          ar1      ma1      ma2      sma1
```

```
##          0.9147  0.8401  0.0945  0.0853
```

```
## s.e.    0.0165  0.0421  0.0409  0.0382
```

```
##
```


```
## sigma^2 estimated as 1.825e-06: log likelihood=3600.53
```

```
## AIC=-7191.05   AICc=-7190.97   BIC=-7168.34
```

```
seas_fcast <- forecast(fit_w_seasonality, h=30)
```

```
plot(seas_fcast)
```


### Forecasts from ARIMA(1,1,2)(0,0,1)[30]



```
# ARIMA
Box.test(fcast_no_holdout$residuals, lag=20, type = "Ljung-Box")

##
## Box-Ljung test
##
## data: fcast_no_holdout$residuals
## X-squared = 15.917, df = 20, p-value = 0.7218
hist(fcast_no_holdout$residuals, breaks = 10)
```

**Histogram of fcast\_no\_holdout\$residuals**




## HoltWinters

```
hw = HoltWinters(deseasonal_humid)
```


```
plot(hw,xaxt='n')
```

## Holt-Winters filtering



```
hw.forecast = forecast(hw, h = 24)  
plot(hw.forecast, xaxt='n')
```

## Forecasts from HoltWinters




```
hw$SSE  
## [1] 0.011527
```

## Testing the distribution of errors in your Holt-Winters model.

```
acf(as.numeric(na.omit(hw.forecast$residuals)),lag.max = 20)
```

**Series** `as.numeric(na.omit(hw.forecast$residuals`



```
# Holtwinters  
Box.test(hw.forecast$residuals,lag=20,type = "Ljung-Box")  
  
##  
## Box-Ljung test  
##  
## data: hw.forecast$residuals  
## X-squared = 203.18, df = 20, p-value < 2.2e-16  
  
hist(hw.forecast$residuals,breaks = 10)
```

**Histogram of hw.forecast\$residuals**

